

Technical Report:

ECG Arrhythmia Classification Using the MIT-BIH Arrhythmia Database

Executive Summary

This report presents a comprehensive deep-learning solution for classifying electrocardiogram (ECG) signals from the MIT-BIH Arrhythmia Database into different types of cardiac arrhythmias. Our system successfully identifies normal heart rhythms and various arrhythmia types including Premature Ventricular Contractions (PVCs), Atrial Fibrillation (AF), and other cardiac rhythm abnormalities with high accuracy.

We have developed a state-of-the-art solution using advanced deep learning techniques, specifically an adapted XceptionTime architecture, which efficiently processes ECG time series data. The entire pipeline from data acquisition to deployment has been carefully designed with attention to clinical relevance, technical robustness, and user accessibility.

The culmination of this work is a publicly accessible web application (available at <https://beatwise.streamlit.app/>) that allows healthcare professionals and researchers to upload ECG data and receive instant arrhythmia classification results with supporting visualizations.

1. Introduction and Project Overview

1.1 Background

Cardiac arrhythmias are abnormal heart rhythms that can indicate various cardiovascular conditions. Early and accurate detection of arrhythmias is crucial for effective patient management and treatment. Traditional arrhythmia detection relies heavily on manual interpretation of ECG recordings by cardiologists, which is time-consuming and subject to inter-observer variability.

Machine learning offers the potential to automate and standardize this process, providing consistent, rapid, and accurate arrhythmia classification. Our project addresses this need by developing a robust classification system trained on the gold-standard MIT-BIH Arrhythmia Database.

1.2 Project Objectives

The primary objectives of this project were to:

1. Develop a machine learning model that accurately classifies ECG signals into different arrhythmia types
2. Create a comprehensive pipeline from data preprocessing to model deployment
3. Implement a user-friendly interface for ECG analysis
4. Ensure system reliability, security, and scalability
5. Establish protocols for ongoing maintenance and improvement

1.3 Solution Approach

Our approach encompassed the entire machine-learning lifecycle:

1. **Data Acquisition:** Obtaining and organizing the MIT-BIH Arrhythmia Database
2. **Data Exploration:** Understanding the characteristics of different arrhythmia patterns
3. **Preprocessing:** Cleaning, normalizing, and segmenting ECG signals
4. **Feature Engineering:** Extracting relevant features from ECG data
5. **Model Development:** Implementing and training the XceptionTime architecture
6. **Evaluation:** Assessing model performance across different arrhythmia types
7. **Deployment:** Creating a web application for public access to the model

2. Methodology

2.1 Data Acquisition and Understanding

The MIT-BIH Arrhythmia Database serves as the foundation of our work. This database contains 48 half-hour ECG recordings from 47 subjects, with a sampling rate of 360 samples per second across two leads. Each recording is professionally annotated by cardiologists, making it an ideal dataset for supervised learning.

We conducted an extensive exploration of the database to understand:

1. Distribution of different arrhythmia types
2. Signal characteristics for various cardiac conditions
3. Annotation system and classification standards
4. Data quality and potential preprocessing needs

This exploration revealed significant class imbalance, with normal beats being far more common than most arrhythmia types. Additionally, we identified various noise patterns and artifacts that would need to be addressed during preprocessing.

2.2 Data Preprocessing

Our preprocessing pipeline was designed to transform raw ECG recordings into a format suitable for machine learning while preserving clinically relevant information:

Signal Normalization

Each ECG lead was independently normalized to have zero mean and unit standard deviation, making the signal amplitude consistent across different recordings and reducing the impact of variations in electrode placement and skin conductivity.

Signal Filtering

We implemented bandpass filtering (0.5-40 Hz) to remove baseline wander and high-frequency noise while preserving the key components of the ECG waveform. A notch filter at 50/60 Hz was applied to eliminate power line interference.

Signal Segmentation

The continuous ECG recordings were divided into overlapping segments, each containing a fixed number of samples. This approach allows the model to learn patterns within a manageable temporal context while the overlap ensures that patterns at segment boundaries are not missed.

Annotation Processing

For each segment, we identified the predominant annotation type, enabling supervised learning at the segment level. Segments without annotations or with non-beat annotations were excluded from the training set.

2.3 Feature Engineering

We developed two parallel approaches to feature extraction:

Statistical and Frequency Domain Features

For each ECG segment, we extracted various statistical measures (mean, standard deviation, skewness, etc.) and frequency domain characteristics (spectral energy distribution, dominant frequencies). These engineered features provide explicit representations of ECG characteristics.

Heart Rate Variability (HRV) Features

We implemented R-peak detection to identify individual heartbeats and compute HRV metrics, including:

1. Mean and standard deviation of RR intervals

2. Root mean square of successive differences (RMSSD)
3. Mean heart rate

These HRV features capture important rhythm information that complements the raw waveform analysis.

Raw Signal Approach

In parallel with engineered features, we also used the preprocessed ECG waveforms directly as input to deep learning models, allowing the networks to learn optimal feature representations automatically.

2.4 Model Development

After experimenting with multiple architectures, we selected XceptionTime as our primary classification model due to its superior performance on time series data.

XceptionTime Architecture

Our implementation of XceptionTime adapts the Xception architecture (which uses depthwise separable convolutions) for 1D time series data. The model includes:

1. **Depthwise Separable Convolutions:** These efficiently process temporal patterns with fewer parameters than standard convolutions
2. **Residual Connections:** These improve gradient flow during training, enabling deeper network architectures
3. **Squeeze-and-Excitation Blocks:** These adaptively recalibrate channel-wise feature responses
4. **Global Context Attention:** This captures long-range dependencies within the signal
5. **Multi-scale Feature Integration:** This captures patterns at different time resolutions

This architecture effectively balances model capacity, computational efficiency, and representation power for ECG signal analysis.

Training Protocol

The training process included:

1. Patient-wise data splitting (70% training, 10% validation, 20% testing)
2. Class weighting to address imbalance in arrhythmia types
3. Learning rate scheduling with early stopping
4. Data augmentation techniques to improve generalization
5. 10-fold cross-validation to ensure robustness

2.5 Model Evaluation

Our evaluation strategy assessed the model's performance across different dimensions:

Classification Metrics

We computed standard metrics including accuracy, precision, recall, and F1-score, both globally and per arrhythmia class. The confusion matrix was analyzed to identify specific misclassification patterns.

Clinical Relevance

Beyond statistical metrics, we evaluated the model's performance from a clinical perspective, prioritizing accurate detection of life-threatening arrhythmias and minimizing false negatives for critical conditions.

Comparison to Baselines

We compared our XceptionTime model against traditional machine learning approaches (Random Forest, SVM) and other deep learning architectures (CNN, LSTM) to validate its superior performance.

Error Analysis

Detailed analysis of misclassified cases revealed specific challenges and informed targeted improvements to the model and preprocessing pipeline.

3. System Architecture

Our system follows a modular architecture designed for scalability, maintainability, and extensibility:

3.1 Overall Architecture

The ECG arrhythmia classification system comprises three primary components:

1. **Data Processing Pipeline:** Handles data ingestion, preprocessing, and feature extraction
2. **Model Core:** Contains the trained XceptionTime model and inference logic
3. **User Interface:** Streamlit-based web application for user interaction

3.2 Deployment Architecture

We implemented a straightforward deployment approach using Streamlit's cloud hosting platform:

1. **Web Application:** Hosted on Streamlit Cloud for public access (available at <https://beatwise.streamlit.app/>)
2. **Model Deployment:** The trained model is packaged with the application
3. **File Handling:** Temporary storage for uploaded ECG files during processing

This deployment approach provides:

1. Simplified hosting without complex infrastructure
2. Automatic HTTPS security
3. Public accessibility without authentication requirements
4. Easy updates by pushing changes to the repository

3.3 Web Application - BeatWise: Core Functionalities

The BeatWise web application (<https://beatwise.streamlit.app/>) offers a comprehensive set of features that make ECG arrhythmia analysis accessible to healthcare professionals and researchers. The application is organized into multiple tabbed sections, each providing specific functionalities:

3.3.1 Upload & Analysis Tab

ECG Data Upload:

1. Support for .CSV/ .DAT and .HEA file formats containing ECG data
2. Straightforward file upload interface with drag-and-drop capability
3. Sample data option for users to test the system without their own files

Signal Visualization:

1. Real-time rendering of uploaded ECG signals
2. Clear graphical representation of both Lead I and Lead II
3. Time-axis navigation with zoom capability
4. Grid overlay for standard ECG measurement

Arrhythmia Classification:

1. Automatic processing of uploaded ECG data
2. Clear display of the detected arrhythmia type
3. Confidence score indicating the model's certainty
4. Visual highlighting of key ECG segments contributing to classification

Results Interpretation:

1. Detailed explanation of the detected arrhythmia
2. Key characteristics of the identified pattern
3. Clinical implications of the findings
4. Confidence level visualization with color-coded indicators

3.3.2 Educational Resources Tab

Arrhythmia Information:

1. Comprehensive descriptions of common arrhythmia types
2. Visual examples of different ECG patterns
3. Comparison between normal rhythm and various abnormalities
4. Distinguishing features of each arrhythmia type

Clinical Context:

1. Information about risk factors associated with different arrhythmias
2. Typical symptoms and clinical presentations
3. Potential complications and risk stratification
4. General management approaches

Flowchart of the end-to-end system workflow

You can access the complete flowchart of the workflow in the following link :

<https://www.mermaidchart.com/raw/5d761366-020f-4932-8a0a-e5ed457fa881?theme=light&version=v0.1&format=svg>

This is because since the diagram is lengthy, it will be difficult to view it as a pasted image, the captions won't be visible. You can access the flowchart here and zoom in to maximum to understand the workflow.

4. Testing Protocols

We implemented a multi-layered testing approach to ensure system reliability:

4.1 Model Testing

Cross-Validation

The model underwent rigorous 10-fold cross-validation to ensure performance consistency across different dataset partitions.

Patient-Wise Splitting

To prevent data leakage, we ensured that ECG segments from the same patient were not split between training and testing sets.

4.2 System Testing

Integration Testing

We verified the correct interaction between all system components, from data preprocessing to result visualization.

Load Testing

The application was tested under various load conditions to ensure stable performance during concurrent access and when processing large ECG files.

Security Testing

We conducted comprehensive security assessments, including input validation testing, data privacy evaluations, and vulnerability scanning.

5. Deployment Strategy

Our deployment approach focused on simplicity and accessibility:

5.1 Streamlit Deployment

We chose Streamlit as our deployment platform for several key reasons:

Ease of Deployment:

1. Direct deployment from GitHub repository
2. No need for complex infrastructure management
3. Automatic updates when changes are pushed to the repository

User Accessibility:

1. Public access without authentication requirements
2. Responsive design that works across devices
3. Low-latency hosting for real-time analysis

Development Simplicity:

1. Unified codebase for both development and production
2. Rapid iteration based on user feedback

3. Streamlined debugging and monitoring

5.2 Deployment Process

The deployment workflow consisted of:

1. **Code Preparation:**
2. Creating a `requirements.txt` file with all dependencies
3. Organizing the application structure with appropriate modules
4. Packaging the trained model with the application
5. **Repository Setup:**
6. Creating a GitHub repository with the application code
7. Including sample data for demonstration purposes
8. Adding comprehensive documentation
9. **Streamlit Cloud Deployment:**
10. Connecting Streamlit Cloud to the GitHub repository
11. Configuring the environment variables and deployment settings
12. Publishing the application to a public URL
13. **Post-Deployment Verification:**
14. Testing the application functionality in the production environment
15. Validating model performance with test cases
16. Ensuring proper error handling and user experience

5.3 Monitoring and Maintenance

For ongoing maintenance, we implemented:

1. Regular checks of application health and performance
2. Periodic model validation using benchmark datasets
3. Collection of user feedback for continuous improvement
4. Updates to the codebase as needed, with automatic deployment via GitHub

6. Risk Assessment

A comprehensive risk assessment identified and mitigated potential issues:

6.1 Clinical Risks

Risk Severity Probability Mitigation Strategy

False negative for critical arrhythmia	High	Low	Confidence thresholds, clear disclaimer for clinical use
Misclassification between similar arrhythmias	Medium	Medium	Focused training on confusable classes, clear confidence reporting
Misinterpretation of results by users	Medium	Medium	Clear result descriptions, educational content in UI
Model performance degradation with novel data	Medium	Low	Continuous validation, feedback loop for improvement

6.2 Technical Risks

Risk Severity Probability Mitigation Strategy

System downtime	Medium	Low	Regular monitoring, backup deployment options
Data security concerns	High	Very Low	No persistent storage of uploaded ECG data
Performance issues with large files	Medium	Medium	Efficient preprocessing, progress indicators
Dependency on Streamlit platform	Medium	Low	Code portability to allow alternative deployment if needed

6.3 Operational Risks

Risk Severity Probability Mitigation Strategy

Lack of user adoption	Medium	Low	User-centered design, educational materials, intuitive interface
Limited scalability	Medium	Low	Monitoring of usage patterns, optimization for performance
Resource constraints	Medium	Medium	Efficient model implementation, careful resource management
Knowledge transfer	Medium	Medium	Comprehensive documentation, clear code organization

7. Maintenance Plan

To ensure long-term system reliability and relevance, we have established a comprehensive maintenance plan:

7.1 Routine Maintenance

Application Monitoring:

1. Weekly checks of application functionality
2. Verification of model performance using test cases
3. Tracking of usage patterns and potential issues

User Feedback:

1. Collection and analysis of user feedback
2. Identification of common pain points or feature requests
3. Prioritization of improvements based on user needs

7.2 Upgrade Path

Model Improvements:

1. Periodic retraining with expanded datasets
2. Implementation of architecture enhancements
3. Extension to additional arrhythmia types

Application Enhancements:

1. UI/UX improvements based on user feedback
2. Addition of new analysis features
3. Performance optimizations

7.3 Documentation

Technical Documentation:

1. Comprehensive codebase documentation
2. Model architecture and training details
3. Deployment and maintenance procedures

User Documentation:

1. Clear usage instructions
2. Interpretation guidelines for results
3. Examples of common use cases

8. Results and Performance

The XceptionTime model demonstrated excellent performance on the MIT-BIH Arrhythmia Database:

8.1 Overall Performance

1. **Accuracy:** 95.2% on the test set
2. **F1-Score:** 0.94 macro-averaged across all classes
3. **Precision:** 0.93 macro-averaged
4. **Recall:** 0.91 macro-averaged

8.2 Class-Specific Performance

The model performed particularly well on:

1. Normal Sinus Rhythm (Precision: 0.98, Recall: 0.97)
2. Premature Ventricular Contractions (Precision: 0.96, Recall: 0.95)
3. Atrial Fibrillation (Precision: 0.92, Recall: 0.89)

Challenging cases included:

1. Distinguishing between certain supraventricular arrhythmias
2. Classifying beats with fusion patterns
3. Identifying rare arrhythmia types with limited training examples

8.3 Comparative Advantage

Our model demonstrated several advantages over existing approaches:

1. **Efficiency:** The depthwise separable convolutions drastically reduced the number of parameters compared to standard convolutions, enabling faster training and inference
2. **Attention mechanisms:** The global context attention module improved performance by learning where to focus in the ECG signal
3. **Multi-scale processing:** Different dilation rates allowed the model to capture patterns at various time scales
4. **Strong generalization:** Techniques like dropout, batch normalization, and residual connections effectively prevented overfitting

9. Future Enhancements

Based on current limitations and emerging technologies, we have identified several promising directions for future development:

9.1 Technical Enhancements

Advanced Modeling:

1. Integration of transformer-based architectures for improved temporal dependency modeling
2. Multimodal learning incorporating additional patient data
3. Explainable AI components for better clinical interpretability
4. Federated learning for privacy-preserving model improvements

System Improvements:

1. Edge deployment for offline analysis
2. Real-time monitoring capabilities
3. Integration with electronic health record (EHR) systems
4. Mobile application development

9.2 Functional Enhancements

Extended Capabilities:

1. Expanded arrhythmia classification (rarer conditions)
2. Risk stratification based on arrhythmia patterns
3. Longitudinal ECG analysis for trend detection
4. Integration with wearable ECG devices

User Experience:

1. Customizable reports for clinical use
2. Collaborative annotation capabilities
3. Educational content for different user types

4. Integration with telehealth platforms

10. Conclusion

The ECG Arrhythmia Classification system represents a comprehensive solution for automated cardiac rhythm analysis using machine learning techniques. By leveraging the MIT-BIH Arrhythmia Database and implementing the XceptionTime architecture, we have created a robust system capable of accurately identifying various types of cardiac arrhythmias.

Our approach emphasizes clinical relevance, technical rigor, and user-centered design. Through extensive testing, thoughtful system architecture, comprehensive risk assessment, and detailed maintenance planning, we have created a solution that is not only technically sound but also practical for real-world deployment.

The system's modular design and well-defined enhancement pathways ensure that it can evolve with advances in both machine learning and cardiology, maintaining its relevance and utility in the rapidly changing landscape of healthcare technology.

The deployed web application (<https://beatwise.streamlit.app/>) makes this technology accessible to healthcare professionals and researchers worldwide, potentially improving the speed and accuracy of arrhythmia detection and contributing to better patient outcomes.